



## Configuration

API Connected

## Query Settings

Documents to retrieve

3

## Upload Document

Choose a PDF file

Drag and drop file here

Limit 200MB per file • PDF

Browse files

## Document Manager

Refresh document list

Total chunks in DB: 0

No documents found in the collection yet.

# RAG Chatbot

Ask questions about your documents and get AI-powered answers with citations

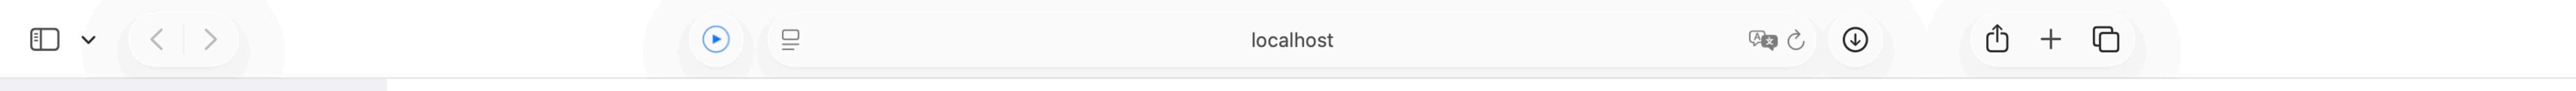
Welcome! Get started by:

1 Upload a PDF Use the sidebar to upload your document

2 Ask Questions Type your question in the chat below

3 Get Answers Receive AI-generated answers with sources

Ask a question about your documents...



### Upload Document

Choose a PDF file

Drag and drop file here  
Limit 200MB per file • PDF

Browse files

app\_buid\_pa... X  
461.7KB

Process Document

---

### Document Manager

Refresh document list

Total chunks in DB: 34

**app\_buid\_paper.pdf**  
Chunks: 34  
 Delete  
app\_buid\_paper.pdf

---

### Danger Zone

I understand this will delete all documents

Clear ALL documents

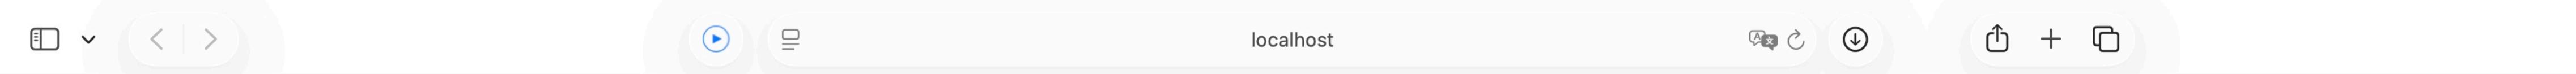
# 🕒 RAG Chatbot

Ask questions about your documents and get AI-powered answers with citations

## Welcome! Get started by:

- 1 **Upload a PDF** Use the sidebar to upload your document
- 2 **Ask Questions** Type your question in the chat below
- 3 **Get Answers** Receive AI-generated answers with sources

Ask a question about your documents... ➤



## Document Manager

Refresh document list

Total chunks in DB: 60

**LoRA.pdf**

Chunks: 26

Delete LoRA.pdf

**app\_buid\_paper.pdf**

Chunks: 34

 Delete  
app\_buid\_paper.pdf

## Danger Zone

 I understand this will delete all documents

Clear ALL documents

## Developed by

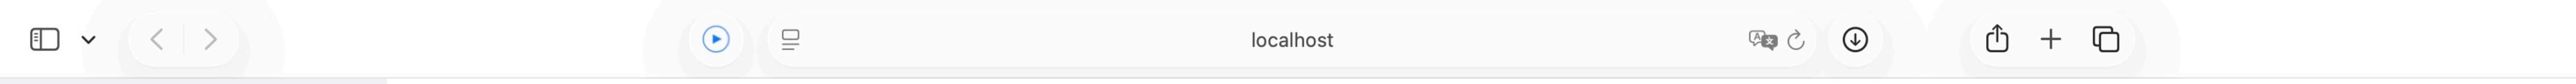
**Pranav Sah** Built with FastAPI + Streamlit +  
Ollama **RAG Chatbot**

Ask questions about your documents and get AI-powered answers with citations

## Welcome! Get started by:

**1 Upload a PDF** Use the sidebar to upload your document**2 Ask Questions** Type your question in the chat below**3 Get Answers** Receive AI-generated answers with sources

Ask a question about your documents...



**Document Manager**

Refresh document list

Total chunks in DB: 60

**LoRA.pdf**  
Chunks: 26  
 Delete LoRA.pdf

**app\_buid\_paper.pdf**  
Chunks: 34  
 Delete app\_buid\_paper.pdf

---

**⚠ Danger Zone**

I understand this will delete all documents  
 Clear ALL documents

---

**Developed by**  
Pranav Sah

Built with FastAPI + Streamlit + Ollama

# RAG Chatbot

Ask questions about your documents and get AI-powered answers with citations

## Welcome! Get started by:

- 1 Upload a PDF** Use the sidebar to upload your document
- 2 Ask Questions** Type your question in the chat below
- 3 Get Answers** Receive AI-generated answers with sources

What is environment scaffolding?

Thinking...

Ask a question about your documents...



## Query Settings

Documents to retrieve

3

## Upload Document

Choose a PDF file

Drag and drop file here

Limit 200MB per file • PDF

[Browse files](#)

## Document Manager

[Refresh document list](#)

Total chunks in DB: 60

### LoRA.pdf

Chunks: 26

[Delete LoRA.pdf](#)

### app\_buid\_paper.pdf

Chunks: 34

[Delete](#)

app\_buid\_paper.pdf

# RAG Chatbot

Ask questions about your documents and get AI-powered answers with citations

## Welcome! Get started by:

1 Upload a PDF Use the sidebar to upload your document

2 Ask Questions Type your question in the chat below

3 Get Answers Receive AI-generated answers with sources

What is environment scaffolding?

According to Document 1 (Page 2), we define environment scaffolding (ES) as:

"...an environment-first paradigm for LLM-based code generation where the model operates inside a structured sandbox that constrains actions and provides continuous, deterministic feedback."

[View 3 Sources](#)

Response time: 39.74s

Ask a question about your documents...

ConfigurationAPI Connected

Query Settings

Documents to retrieve 3

Upload Document

Choose a PDF file  
Drag and drop file here  
Limit 200MB per file • PDF  
[Browse files](#)

Document Manager

Total chunks in DB: 60

<b>LoRA.pdf</b>	Chunks: 26	<a href="#">Delete LoRA.pdf</a>
<b>app_buid_paper.pdf</b>	Chunks: 34	<a href="#">Delete app_buid_paper.pdf</a>

## Welcome! Get started by:

- 1 Upload a PDF** Use the sidebar to upload your document
- 2 Ask Questions** Type your question in the chat below
- 3 Get Answers** Receive AI-generated answers with sources

What is environment scaffolding?

According to Document 1 (Page 2), we define environment scaffolding (ES) as:  
"...an environment-first paradigm for LLM-based code generation where the model operates inside a structured sandbox that constrains actions and provides continuous, deterministic feedback."

[View 3 Sources](#)

**1. app\_buid\_paper.pdf (Page 2)**

Figure 1: Environment scaffolding vs. model-centric generation.  
ES wraps the model with a finite, validated workflow that catches errors early and repairs them before proceeding. either training or prompt engineering, but this paradigm fails to address fundamental reliability issues inherent in probabilistic generation. Recent surveys [Jiang et al., 2024, Paul et al., 2024] note the field requires a shift from model-centric to environment-cen...

**2. app\_buid\_paper.pdf (Page 3)**

Table 1: Environment scaffolding (ES) vs. model-centric generation.  
Aspect Model-Centric Environment Scaffolding (Ours) Task decomposition Single/loosely guided multi-step; no fixed structure Explicit pipeline (FSM): schema → API → UI Validation Late or ad-hoc checks Integration

**3. app\_buid\_paper.pdf (Page 5)**

Figure 2: app.build architecture expressed through environment scaffolding.  
The orchestrator plans stages per stack; each sub-task runs in a sandbox, is validated, and only then merged. CI/CD and DB provisioning are integrated.  
4.1 Evaluation Framework 4.2 Prompt Dataset The evaluation dataset comprises 30 prompts designed to assess system performance across diverse application development scenarios. Independent human contributors with no pri...

Ask a question about your documents...

Danger Zone

The screenshot shows a web application interface for a RAG Chatbot. The top navigation bar includes standard browser controls (back, forward, search, etc.) and a title 'localhost'. On the far right are 'Deploy' and a three-dot menu icon.

**API Connected** (green status indicator)

**Query Settings**

Documents to retrieve: 3 (indicated by a red slider)

**Upload Document**

Choose a PDF file

Drag and drop file here  
Limit 200MB per file • PDF

Browse files

**Document Manager**

Refresh document list

Total chunks in DB: 60

**LoRA.pdf**

Chunks: 26

Delete LoRA.pdf

**app\_buid\_paper.pdf**

Chunks: 34

Delete app\_buid\_paper.pdf

**RAG Chatbot**

Ask questions about your documents and get AI-powered answers with citations

What is environment scaffolding?

According to Document 1, Page 2:  
"We define environment scaffolding (ES) as an environment-first paradigm for LLM-based code generation where the model operates inside a structured sandbox that constrains actions and provides continuous, deterministic feedback."

Additionally, from Document 1, Page 2:  
"Rather than relying on larger models or prompt-only techniques, ES improves the context around the model — shaping the action space, providing templates and tools, and validating each step — so that creativity is channeled into safe, verifiable outcomes."

> View 3 Sources

Response time: 56.32s

What is LoRA in large language models?

According to Document 3 - Source: LoRA.pdf, Page: 11:  
"LoRA achieves comparable performance to full fine-tuning while drastically reducing computational costs."  
However, the definition of LoRA itself is not explicitly stated. But based on the context and related information in the document, we can infer that LoRA stands for "Low-Rank Adaptation", which is a method used for adapting large language models to downstream tasks with reduced computational costs.  
A more detailed explanation is provided in Document 2 - Source: LoRA.pdf, Page: 9:  
"By freezing the majority of the model's parameters and training only small, low-rank matrices, LoRA achieves comparable performance to full fine-tuning while drastically reducing computational costs."

> View 3 Sources

Response time: 58.33s

Ask a question about your documents...



API Connected

What is LoRA in large language models?

According to Document 3 - Source: LoRA.pdf, Page: 11:

"LoRA achieves comparable performance to full fine-tuning while drastically reducing computational costs."

However, the definition of LoRA itself is not explicitly stated. But based on the context and related information in the document, we can infer that LoRA stands for "Low-Rank Adaptation", which is a method used for adapting large language models to downstream tasks with reduced computational costs.

A more detailed explanation is provided in Document 2 - Source: LoRA.pdf, Page: 9:

"By freezing the majority of the model's parameters and training only small, low-rank matrices, LoRA achieves comparable performance to full fine-tuning while drastically reducing computational costs."

## Query Settings

Documents to retrieve

3

## Upload Document

Choose a PDF file

Drag and drop file here

Limit 200MB per file • PDF

Browse files

## Document Manager

Refresh document list

Total chunks in DB: 60

### LoRA.pdf

Chunks: 26

Delete LoRA.pdf

### app\_buid\_paper.pdf

Chunks: 34

Delete

app\_buid\_paper.pdf

## Danger Zone

 I understand this will delete all documents

View 3 Sources

#### 1. app\_buid\_paper.pdf (Page 10)

Quoc Le, and Charles Sutton.



Program synthesis with large language models, 2021.

URL <https://arxiv.org/abs/2108.07732>.

Valerii Babushkin and Arseny Kravchenko.

Machine Learning System Design with End-to-End Examples.

Manning Publications, 2025.

Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Mic...

#### 2. LoRA.pdf (Page 9)

6 Related Works 6.1 Transformer Language Models The Transformer architecture, as introduced by Vaswani et al.

(2017), has proven to be a highly effective sequence-to-sequence model due to its heavy use of self-attention mechanisms.

Radford et al.

(2018) applied it to autoregressive language modeling, significantly boosting its utility in the field.

Since then, Transformer-based models have become a staple in natural language processing (NLP),...

#### 3. LoRA.pdf (Page 11)

pting large language models for downstream tasks.

By freezing the majority of the model's parameters and training only small, low-rank matrices, LoRA achieves comparable performance to full fine-tuning while drastically reducing computational costs.

Ask a question about your documents...



Clear ALL documents